R Workshop Featuring dplyr

Jorge Loría

October 20th, 2020

What is `dplyr`?????

What is `dplyr`?????

Dplyr is a package that permits a *nice* interface in R to work with data.frames.

_____

What is `dplyr`?????

Dplyr is a package that permits a *nice* interface in R to work with data.frames.

Ok, nice, but what is a data.frame?

---

[1] categorical
[2] https://vita.had.co.nz/papers/tidy-data.pdf

What is `dplyr`?????

Dplyr is a package that permits a *nice* interface in R to work with data.frames.

Ok, nice, but what is a data.frame?

A data.frame is like a matrix of dimensions $n \times p$, where we have several different types of data. Each column corresponds to a single variable, and each variable has a specific type (numeric, string, logical, factor[1]). Each row should correspond to a single observation[2].

---

[1]categorical
[2]https://vita.had.co.nz/papers/tidy-data.pdf

What dataset are we going to use?

We are going to use a dataset from `tidytuesday`, about Himalayan Climbers

What dataset are we going to use?

We are going to use a dataset from `tidytuesday`, about Himalayan Climbers

Figure 1: Himalaya? It must be easy to survive there

How to access it?

```r
install.packages('tidytuesdayR')
library(tidytuesdayR)
himalaya <- tidytuesdayR::tt_load('2020-09-22')
members <- himalaya$members
```

Let's start with questions

What questions can we ask this dataset that you think would be interesting to know? (Please write them in the chat!)

Let's start with questions

What questions can we ask this dataset that you think would be interesting to know? (Please write them in the chat!)

The idea is to be able to solve most, if not all of those questions by the end of the workshop!

Let's start with questions

What questions can we ask this dataset that you think would be interesting to know? (Please write them in the chat!)

The idea is to be able to solve most, if not all of those questions by the end of the workshop!

We will study some tools, and you should be able to answer them by the end of the workshop. If I can see that's not possible, I'll tell you what libraries are good to answer them. Ok?

Let's take a look:

```
members %>%
  head()

## # A tibble: 6 x 21
##   expedition_id member_id  peak_id peak_name  year season sex     age ci
##   <chr>         <chr>      <chr>   <chr>     <dbl> <chr>  <chr> <dbl> <cl
## 1 AMAD78301     AMAD7830~  AMAD    Ama Dabl~  1978 Autumn M        40 Fra
## 2 AMAD78301     AMAD7830~  AMAD    Ama Dabl~  1978 Autumn M        41 Fra
## 3 AMAD78301     AMAD7830~  AMAD    Ama Dabl~  1978 Autumn M        27 Fra
## 4 AMAD78301     AMAD7830~  AMAD    Ama Dabl~  1978 Autumn M        40 Fra
## 5 AMAD78301     AMAD7830~  AMAD    Ama Dabl~  1978 Autumn M        34 Fra
## 6 AMAD78301     AMAD7830~  AMAD    Ama Dabl~  1978 Autumn M        25 Fra
## # ... with 12 more variables: expedition_role <chr>, hired <lgl>,
## #   highpoint_metres <dbl>, success <lgl>, solo <lgl>, oxygen_used <lgl
## #   died <lgl>, death_cause <chr>, death_height_metres <dbl>, injured <l
## #   injury_type <chr>, injury_height_metres <dbl>
```

%>%

This is a function that doesn't do much, but does a lot. It allows to compose functions (as the math people do) but in a way that permits an easy reading of the functions, and what is happening.

%>%

This is a function that doesn't do much, but does a lot. It allows to compose functions (as the math people do) but in a way that permits an easy reading of the functions, and what is happening. So, instead of writing: $f(g(h(i(x))))$

%>%

This is a function that doesn't do much, but does a lot. It allows to compose functions (as the math people do) but in a way that permits an easy reading of the functions, and what is happening. So, instead of writing: $f(g(h(i(x))))$, we write:

```
x %>%
  i() %>%
  h() %>%
  g() %>%
  f()
```

Which is arguably easier to read than the previous expression. Specially if some of those functions had extra arguments.

The Verbs

There are 6 main verbs in `dplyr` that we will study:

function	action
filter	keeps rows that satisfy a condition
arrange	sorts the rows following the order
select	keeps/eliminates the columns by name
mutate	creates new variables from existing variables
summarise	summarises the data
groub_by*	groups under specific conditions

The Verbs

All the verbs work very similarly (in `dplyr`, in English verbs are more confusing):

The Verbs

All the verbs work very similarly (in dplyr, in English verbs are more confusing):

- first argument is a dataframe,

The Verbs

All the verbs work very similarly (in `dplyr`, in English verbs are more confusing):

▶ first argument is a dataframe,

▶ arguments describe what to do with the dataframe, using the existing variables (without quotations).

The Verbs

All the verbs work very similarly (in `dplyr`, in English verbs are more confusing):

▶ first argument is a dataframe,

▶ arguments describe what to do with the dataframe, using the existing variables (without quotations).

▶ output is a dataframe

The Verbs

All the verbs work very similarly (in `dplyr`, in English verbs are more confusing):

- ▶ first argument is a dataframe,

- ▶ arguments describe what to do with the dataframe, using the existing variables (without quotations).

- ▶ output is a dataframe

The Verbs

All the verbs work very similarly (in `dplyr`, in English verbs are more confusing):

▶ first argument is a dataframe,

▶ arguments describe what to do with the dataframe, using the existing variables (without quotations).

▶ output is a dataframe

This structure allows us to concatenate (`%>%`) simple operations to obtain complex results.

filter

Figure 2: Another type of filters

filter

This functions removes rows that don't satisfy a (or several) condition that we specify.
The arguments it receives are logical, and will use it to do that removal:

```r
library(dplyr) ## load the library
members %>%
  filter(oxygen_used)
```

```
## # A tibble: 18,233 x 21
##   expedition_id member_id peak_id peak_name  year season sex   age
##   <chr>         <chr>     <chr>   <chr>     <dbl> <chr>  <chr> <dbl>
## 1 ANN170101     ANN17010~ ANN1    Annapurn~  1970 Spring M        35
## 2 ANN170102     ANN17010~ ANN1    Annapurn~  1970 Spring M        28
## 3 ANN170102     ANN17010~ ANN1    Annapurn~  1970 Spring M        32
## 4 ANN177301     ANN17730~ ANN1    Annapurn~  1977 Autumn M        37
## 5 ANN177301     ANN17730~ ANN1    Annapurn~  1977 Autumn M        29
## 6 ANN177301     ANN17730~ ANN1    Annapurn~  1977 Autumn M        28
## 7 ANN178301     ANN17830~ ANN1    Annapurn~  1978 Autumn F        35
```

filter

We can use several columns to filter, and can even modify them. Let's see what people older than 75 years **didn't** need to use oxygen

```
members %>%
  filter(!oxygen_used,age > 75)
```

```
## # A tibble: 27 x 21
##    expedition_id member_id peak_id peak_name  year season sex    age
##    <chr>         <chr>     <chr>   <chr>     <dbl> <chr>  <chr> <dbl>
## 1 GIMM93301     GIMM9330~ GIMM    Gimmigel~  1993 Autumn M        80
## 2 GIME93301     GIME9330~ GIME    Gimmigel~  1993 Autumn M        80
## 3 PUMO05104     PUMO0510~ PUMO    Pumori     2005 Spring M        76
## 4 AMAD06347     AMAD0634~ AMAD    Ama Dabl~  2006 Autumn M        76
## 5 AMAD07314     AMAD0731~ AMAD    Ama Dabl~  2007 Autumn M        77
## 6 YAKA08201     YAKA0820~ YAKA    Yakawa K~  2008 Summer F        77
## 7 AMAD08331     AMAD0833~ AMAD    Ama Dabl~  2008 Autumn M        78
## 8 TASH10302     TASH1030~ TASH    Tashi Ka~  2010 Autumn M        76
```

filter

We can also use several columns at once to do a filter. Let's see what climber(s?) died
a little bit after getting injured:

```
members %>%
  filter(death_height_metres > injury_height_metres)
```

```
## # A tibble: 1 x 21
##   expedition_id member_id peak_id peak_name  year season sex    age ci
##   <chr>         <chr>     <chr>   <chr>     <dbl> <chr>  <chr> <dbl> <c
## 1 BAUD70101     BAUD7010~ BAUD    Baudha     1970 Spring M        26 Ja
## # ... with 12 more variables: expedition_role <chr>, hired <lgl>,
## #   highpoint_metres <dbl>, success <lgl>, solo <lgl>, oxygen_used <lgl
## #   died <lgl>, death_cause <chr>, death_height_metres <dbl>, injured <
## #   injury_type <chr>, injury_height_metres <dbl>
```

Exercises `filter`

Exercises `filter`

▶ From those who where not injured, how many died in 1979?

Exercises `filter`

▶ From those who where not injured, how many died in 1979?

▶ How many of those that succeeded, were doing it solo and got injured above 1000 meters?

Exercises `filter`

▶ From those who where not injured, how many died in 1979?

▶ How many of those that succeeded, were doing it solo and got injured above 1000 meters?

▶ How many people died on their climb below 1000 meters?

Exercises `filter`

▶ From those who where not injured, how many died in 1979?

▶ How many of those that succeeded, were doing it solo and got injured above 1000 meters?

▶ How many people died on their climb below 1000 meters?

Exercises `filter`

► From those who where not injured, how many died in 1979?

► How many of those that succeeded, were doing it solo and got injured above 1000 meters?

► How many people died on their climb below 1000 meters?

filter saving

Let's now save the dataframe of the members that went on a expedition in 1905:

```r
members1905 <- members %>%
  filter(year == 1905)
```

arrange

arrange

Figure 3: Another type of arrangement

arrange

arrange

This verb sorts the data frame with the column(s) that we tell it to use

```
members %>%
  arrange(year)
```

```
## # A tibble: 76,519 x 21
##    expedition_id member_id  peak_id peak_name  year season sex    age
##    <chr>         <chr>      <chr>   <chr>     <dbl> <chr>  <chr> <dbl>
## 1 KANG05201      KANG0520~  KANG    Kangchen~  1905 Summer M        29
## 2 KANG05201      KANG0520~  KANG    Kangchen~  1905 Summer M        NA
## 3 KANG05201      KANG0520~  KANG    Kangchen~  1905 Summer M        NA
## 4 KANG05201      KANG0520~  KANG    Kangchen~  1905 Summer M        NA
## 5 KANG05201      KANG0520~  KANG    Kangchen~  1905 Summer M        NA
## 6 KANG05201      KANG0520~  KANG    Kangchen~  1905 Summer M        NA
## 7 KANG05201      KANG0520~  KANG    Kangchen~  1905 Summer M        36
## 8 KANG05201      KANG0520~  KANG    Kangchen~  1905 Summer M        NA
## 9 KANG05201      KANG0520~  KANG    Kangchen~  1905 Summer M        NA
```

arrange on characters

Let's go back to the 1905 dataset, and check how it orders when we use a string, instead of a number:

```
members1905 %>%
  arrange(member_id)
```

```
## # A tibble: 9 x 21
##   expedition_id member_id peak_id peak_name   year season sex     age ci
##   <chr>         <chr>     <chr>   <chr>      <dbl> <chr>  <chr> <dbl> <cl
## 1 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        29 UK
## 2 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        36 Sw
## 3 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        NA Sw
## 4 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        NA Sw
## 5 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        NA It
## 6 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        NA Ne
## 7 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        NA Ne
## 8 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        NA Ne
```

arrange on strings

arrange uses the lexicographic order, to sort when it encounters letters/characters, so: $a < A < b < B \cdots < z < Z$. If you're not sure if a number or a symbol is smaller than another one, you can try it out on the 'Console' in R for example type `'1'<'a'`. Is this true? Or false?

arrange on strings

arrange uses the lexicographic order, to sort when it encounters letters/characters, so: $a < A < b < B \cdots < z < Z$. If you're not sure if a number or a symbol is smaller than another one, you can try it out on the 'Console' in R for example type `'1'<'a'`. Is this true? Or false? TRUE!

arrange on strings

arrange uses the lexicographic order, to sort when it encounters letters/characters, so: $a < A < b < B \cdots < z < Z$. If you're not sure if a number or a symbol is smaller than another one, you can try it out on the 'Console' in R for example type '1'<'a'. Is this true? Or false? TRUE!How about '&' > '2'?

arrange on strings

arrange uses the lexicographic order, to sort when it encounters letters/characters, so: $a < A < b < B \cdots < z < Z$. If you're not sure if a number or a symbol is smaller than another one, you can try it out on the 'Console' in R for example type '1'<'a'. Is this true? Or false? TRUE!How about '&' > '2'?FALSE!

arrange on strings

arrange uses the lexicographic order, to sort when it encounters letters/characters, so: $a < A < b < B \cdots < z < Z$. If you're not sure if a number or a symbol is smaller than another one, you can try it out on the 'Console' in R for example type `'1'<'a'`. Is this true? Or false? TRUE!How about `'&' > '2'`?FALSE! Is `' ' < '1'` (space less than one)?

arrange on strings

arrange uses the lexicographic order, to sort when it encounters letters/characters, so: $a < A < b < B \cdots < z < Z$. If you're not sure if a number or a symbol is smaller than another one, you can try it out on the 'Console' in R for example type `'1'<'a'`. Is this true? Or false? TRUE! How about `'&' > '2'`? FALSE! Is `' ' < '1'` (space less than one)? TRUE!

arrange on several columns

When we include several arguments, it first sorts the first one, then the second one *within* the first order, then the third one *within* the second one, and so on...

```
members %>%
  arrange(desc(year),citizenship,hired,peak_name)
```

```
## # A tibble: 76,519 x 21
##    expedition_id member_id  peak_id peak_name  year season sex     age
##    <chr>         <chr>      <chr>   <chr>     <dbl> <chr>  <chr> <dbl>
##  1 AMAD19107     AMAD1910~  AMAD    Ama Dabl~  2019 Spring M        55
##  2 EVER19116     EVER1911~  EVER    Everest    2019 Spring M        53
##  3 EVER19117     EVER1911~  EVER    Everest    2019 Spring M        38
##  4 EVER19117     EVER1911~  EVER    Everest    2019 Spring M        40
##  5 EVER19148     EVER1914~  EVER    Everest    2019 Spring M        51
##  6 KANG19102     KANG1910~  KANG    Kangchen~  2019 Spring M        43
##  7 EVER19183     EVER1918~  EVER    Everest    2019 Spring M        48
##  8 AMAD19106     AMAD1910~  AMAD    Ama Dabl~  2019 Spring F        33
```

arrange on transformations

First, let's only take those that got injured, and then see who got furthest away with respect to their injury. Is there anything weird going on? Is my code correct? Is the data correct?

```
members %>%
  filter(injured) %>%
  arrange(highpoint_metres - injury_height_metres)

## # A tibble: 1,713 x 21
##    expedition_id member_id peak_id peak_name  year season sex    age
##    <chr>         <chr>     <chr>   <chr>     <dbl> <chr>  <chr> <dbl>
## 1 CHOY02327     CHOY0232~ CHOY    Cho Oyu    2002 Autumn M        29
## 2 MAKA15109     MAKA1510~ MAKA    Makalu     2015 Spring M        50
## 3 MAKA15109     MAKA1510~ MAKA    Makalu     2015 Spring M        46
## 4 CHOY03301     CHOY0330~ CHOY    Cho Oyu    2003 Autumn M        53
## 5 HIML13301     HIML1330~ HIML    Himlung ~  2013 Autumn M        44
## 6 KANG00106     KANG0010~ KANG    Kangchen~  2000 Spring M        35
```

Exercises for `arrange`:

Exercises for `arrange`:

- From those that died, sort by those who got the highest before dying w.r.t. where they died.

Exercises for `arrange`:

- From those that died, sort by those who got the highest before dying w.r.t. where they died.

- Sort by year and then season. Does this make sense? If it doesn't, how could we fix it?

Exercises for `arrange`:

▶ From those that died, sort by those who got the highest before dying w.r.t. where they died.

▶ Sort by year and then season. Does this make sense? If it doesn't, how could we fix it?

▶ If you arrange by `died`, can you tell how `arrange` interprets the logicals?

select

select

This verbs selects the columns that we want to keep. Sometimes, we only need a
couple of variables, and it's good to get rid of the rest:

```
members %>%
  select(year,sex,citizenship)
```

```
## # A tibble: 76,519 x 3
##     year sex   citizenship
##    <dbl> <chr> <chr>
## 1  1978 M     France
## 2  1978 M     France
## 3  1978 M     France
## 4  1978 M     France
## 5  1978 M     France
## 6  1978 M     France
## 7  1978 M     France
## 8  1978 M     France
```

select ranges

Sometimes we want to select all columns between two other columns, so we can use the colon (:, not the organ) to do this:

select ranges

Sometimes we want to select all columns between two other columns, so we can use the colon (:, not the organ) to do this:

```
members %>%
  select(age:success)
```

```
## # A tibble: 76,519 x 6
##       age citizenship expedition_role hired highpoint_metres success
##     <dbl> <chr>       <chr>           <lgl>            <dbl> <lgl>
## 1     40 France      Leader          FALSE               NA FALSE
## 2     41 France      Deputy Leader   FALSE             6000 FALSE
## 3     27 France      Climber         FALSE               NA FALSE
## 4     40 France      Exp Doctor      FALSE             6000 FALSE
## 5     34 France      Climber         FALSE               NA FALSE
## 6     25 France      Climber         FALSE             6000 FALSE
## 7     41 France      Climber         FALSE             6000 FALSE
## 8     29 France      Climber         FALSE             6000 FALSE
```

Eliminating with `select`

Other times we want to get rid of specific variables. For this, we can use the - (minus) symbol.

```
members %>%
  select(-expedition_id,-member_id,-peak_id)
```

```
## # A tibble: 76,519 x 18
##    peak_name  year season sex     age citizenship expedition_role hired
##    <chr>     <dbl> <chr>  <chr> <dbl> <chr>       <chr>           <lgl>
## 1 Ama Dabl~  1978 Autumn M        40 France      Leader          FALSE
## 2 Ama Dabl~  1978 Autumn M        41 France      Deputy Leader   FALSE
## 3 Ama Dabl~  1978 Autumn M        27 France      Climber         FALSE
## 4 Ama Dabl~  1978 Autumn M        40 France      Exp Doctor      FALSE
## 5 Ama Dabl~  1978 Autumn M        34 France      Climber         FALSE
## 6 Ama Dabl~  1978 Autumn M        25 France      Climber         FALSE
## 7 Ama Dabl~  1978 Autumn M        41 France      Climber         FALSE
## 8 Ama Dabl~  1978 Autumn M        29 France      Climber         FALSE
```

More of `select`

We can also use the number of the column to indicate which columns to select, and combine it with the names.

```
members %>%
  select(4,6:10,highpoint_metres)
```

```
## # A tibble: 76,519 x 7
##    peak_name  season sex     age citizenship expedition_role highpoint_r
##    <chr>      <chr>  <chr> <dbl> <chr>       <chr>
## 1 Ama Dablam Autumn M        40 France      Leader
## 2 Ama Dablam Autumn M        41 France      Deputy Leader
## 3 Ama Dablam Autumn M        27 France      Climber
## 4 Ama Dablam Autumn M        40 France      Exp Doctor
## 5 Ama Dablam Autumn M        34 France      Climber
## 6 Ama Dablam Autumn M        25 France      Climber
## 7 Ama Dablam Autumn M        41 France      Climber
## 8 Ama Dablam Autumn M        29 France      Climber
```

More `select` functions:

There's a family of functions designed to work with select, so we can work more easily. Among them we have:

- ▶ `starts_with('some')` that looks for all columns that begin with `'some'`

More `select` functions:

There's a family of functions designed to work with select, so we can work more easily. Among them we have:

▶ `starts_with('some')` that looks for all columns that begin with `'some'`

▶ `ends_with('thing')` that looks for all the columns that start with `'thing'`

More `select` functions:

There's a family of functions designed to work with select, so we can work more easily. Among them we have:

▶ `starts_with('some')` that looks for all columns that begin with `'some'`

▶ `ends_with('thing')` that looks for all the columns that start with `'thing'`

▶ `contains('else')` that looks for all the columns that contain `'else'` in their name

More select functions:

There's a family of functions designed to work with select, so we can work more easily. Among them we have:

- ▶ starts_with('some') that looks for all columns that begin with 'some'

- ▶ ends_with('thing') that looks for all the columns that start with 'thing'

- ▶ contains('else') that looks for all the columns that contain 'else' in their name

- ▶ matches(...) that looks for all the columns that match with a regular expression (see ?regexp) that we indicate.

More select functions:

There's a family of functions designed to work with select, so we can work more easily. Among them we have:

▶ `starts_with('some')` that looks for all columns that begin with `'some'`

▶ `ends_with('thing')` that looks for all the columns that start with `'thing'`

▶ `contains('else')` that looks for all the columns that contain `'else'` in their name

▶ `matches(...)` that looks for all the columns that match with a regular expression (see `?regexp`) that we indicate.

▶ `num_range('x',1:15)` that looks for all columns that are like x1, x2,...,x15

More select functions:

There's a family of functions designed to work with select, so we can work more easily. Among them we have:

- ▶ starts_with('some') that looks for all columns that begin with 'some'

- ▶ ends_with('thing') that looks for all the columns that start with 'thing'

- ▶ contains('else') that looks for all the columns that contain 'else' in their name

- ▶ matches(...) that looks for all the columns that match with a regular expression (see ?regexp) that we indicate.

- ▶ num_range('x',1:15) that looks for all columns that are like x1, x2,...,x15

More select functions:

There's a family of functions designed to work with select, so we can work more easily. Among them we have:

- ▶ starts_with('some') that looks for all columns that begin with 'some'

- ▶ ends_with('thing') that looks for all the columns that start with 'thing'

- ▶ contains('else') that looks for all the columns that contain 'else' in their name

- ▶ matches(...) that looks for all the columns that match with a regular expression (see ?regexp) that we indicate.

- ▶ num_range('x',1:15) that looks for all columns that are like x1, x2,...,x15

Select all the columns that end with 'ed', and the column solo. Is there something that stands out?

More select functions:

There's a family of functions designed to work with select, so we can work more easily. Among them we have:

▶ starts_with('some') that looks for all columns that begin with 'some'

▶ ends_with('thing') that looks for all the columns that start with 'thing'

▶ contains('else') that looks for all the columns that contain 'else' in their name

▶ matches(...) that looks for all the columns that match with a regular expression (see ?regexp) that we indicate.

▶ num_range('x',1:15) that looks for all columns that are like x1, x2,...,x15

Select all the columns that end with 'ed', and the column solo. Is there something that stands out? Define a dataframe named members_n that only has members that died, and no columns end with 'ed' or 'id'.

Renaming with select

We can use select to rename the columns we are choosing:

```
members %>%
  select(gender = sex,used_oxygen = oxygen_used,role = expedition_role)

## # A tibble: 76,519 x 3
##    gender used_oxygen role
##    <chr>  <lgl>       <chr>
##  1 M      FALSE       Leader
##  2 M      FALSE       Deputy Leader
##  3 M      FALSE       Climber
##  4 M      FALSE       Exp Doctor
##  5 M      FALSE       Climber
##  6 M      FALSE       Climber
##  7 M      FALSE       Climber
##  8 M      FALSE       Climber
##  9 M      FALSE       Climber
```

rename

What did you notice from the previous slide?

rename

What did you notice from the previous slide? We lost all of the columns we didn't mentioned. We can modify this by including everything() next to the last column. Or we can use the function rename.

```
members %>%
  rename(gender = sex,used_oxygen = oxygen_used,role = expedition_role)
```

```
## # A tibble: 76,519 x 21
##    expedition_id member_id peak_id peak_name  year season gender   age
##    <chr>         <chr>     <chr>   <chr>     <dbl> <chr>  <chr>  <dbl>
## 1 AMAD78301     AMAD7830~ AMAD    Ama Dabl~  1978 Autumn M         40
## 2 AMAD78301     AMAD7830~ AMAD    Ama Dabl~  1978 Autumn M         41
## 3 AMAD78301     AMAD7830~ AMAD    Ama Dabl~  1978 Autumn M         27
## 4 AMAD78301     AMAD7830~ AMAD    Ama Dabl~  1978 Autumn M         40
## 5 AMAD78301     AMAD7830~ AMAD    Ama Dabl~  1978 Autumn M         34
## 6 AMAD78301     AMAD7830~ AMAD    Ama Dabl~  1978 Autumn M         25
## 7 AMAD78301     AMAD7830~ AMAD    Ama Dabl~  1978 Autumn M         41
```

mutate

`mutate`

Figure 4: Another type of mutate

mutate

Let's concatenate the role and the citizenship, using the function `paste`, which binds together two strings:

mutate

Let's concatenate the role and the citizenship, using the function `paste`, which binds together two strings:

```
members_n %>%
  mutate(Role_citizenship = paste(expedition_role,citizenship))
```

```
## # A tibble: 1,106 x 15
##    peak_name  year season sex     age citizenship expedition_role
##    <chr>     <dbl> <chr>  <chr> <dbl> <chr>       <chr>
##  1 Ama Dabl~  1979 Autumn M        23 New Zealand Climber
##  2 Ama Dabl~  1983 Autumn M        31 Switzerland Leader
##  3 Ama Dabl~  1983 Autumn F        28 Switzerland Climber
##  4 Ama Dabl~  1985 Spring M        32 Japan       Climber
##  5 Ama Dabl~  1988 Spring M        33 Canada      Climber
##  6 Ama Dabl~  1992 Spring M        36 Spain       Leader
##  7 Annapurn~  1970 Spring M        32 UK          Climber
##  8 Annapurn~  1973 Spring M        37 Japan       Climber
```

mutate

We can also do numeric operations, for example getting differences explicitly:

```
members_n %>%
  mutate(Difference_mts_died = death_height_metres - highpoint_metres)
```

```
## # A tibble: 1,106 x 15
##    peak_name  year season sex    age citizenship expedition_role
##    <chr>     <dbl> <chr>  <chr> <dbl> <chr>       <chr>
##  1 Ama Dabl~  1979 Autumn M        23 New Zealand Climber
##  2 Ama Dabl~  1983 Autumn M        31 Switzerland Leader
##  3 Ama Dabl~  1983 Autumn F        28 Switzerland Climber
##  4 Ama Dabl~  1985 Spring M        32 Japan       Climber
##  5 Ama Dabl~  1988 Spring M        33 Canada      Climber
##  6 Ama Dabl~  1992 Spring M        36 Spain       Leader
##  7 Annapurn~  1970 Spring M        32 UK          Climber
##  8 Annapurn~  1973 Spring M        37 Japan       Climber
##  9 Annapurn~  1973 Spring M        36 Japan       Climber
```

mutate to create brand new columns

We can also add and create our own columns, using our own values or from other places. But we have to be very careful they are in the appropriate order, otherwise we risk making a very dangerous mistake. How can it be dangerous?

```
members1905 %>% # this has 9 rows!
  mutate(my_row = (1:9)^2 + log(15)*(9-row_number()))
```

```
## # A tibble: 9 x 22
##   expedition_id member_id peak_id peak_name   year season sex     age ci
##   <chr>         <chr>     <chr>   <chr>      <dbl> <chr>  <chr> <dbl> <cl
## 1 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        29 UK
## 2 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        NA Sw
## 3 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        NA Ne
## 4 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        NA Ne
## 5 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        NA Ne
## 6 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        NA Ne
## 7 KANG05201     KANG0520~ KANG    Kangchen~   1905 Summer M        36 Sw
```

mutate + if_else

R has a very neat function called if_else that is just like an if, and checks wethere a condition is true, then do something, if it's not, then do something else:

```r
members1905 %>% ## notice the difference between: ' and "
  mutate(Cheating = if_else(expedition_role == 'Leader',
                            'Cheated',
                            "Didn't cheat"))
```

```
## # A tibble: 9 x 22
##   expedition_id member_id peak_id peak_name  year season sex     age ci
##   <chr>         <chr>     <chr>   <chr>     <dbl> <chr>  <chr> <dbl> <c
## 1 KANG05201     KANG0520~ KANG    Kangchen~  1905 Summer M        29 UK
## 2 KANG05201     KANG0520~ KANG    Kangchen~  1905 Summer M        NA Sw
## 3 KANG05201     KANG0520~ KANG    Kangchen~  1905 Summer M        NA Ne
## 4 KANG05201     KANG0520~ KANG    Kangchen~  1905 Summer M        NA Ne
## 5 KANG05201     KANG0520~ KANG    Kangchen~  1905 Summer M        NA Ne
## 6 KANG05201     KANG0520~ KANG    Kangchen~  1905 Summer M        NA Ne
```

mutate + *other functions*

We can also combine mutate with other functions to obtain new columns that depend
on all the values from specific columns of the dataframe:

```
members1905 %>%
  mutate(anyone_died = any(died),
         max_height = max(highpoint_metres,na.rm = TRUE),
         last_citizenship = last(citizenship),
         youngest = min(age),
         percent_dead = mean(died),
         diff_from_average_height_death =death_height_metres- mean(death_he
  select(-(1:21)) # to see the new variables
```

```
## # A tibble: 9 x 6
##   anyone_died max_height last_citizenship youngest percent_dead diff_fr
##   <lgl>           <dbl> <chr>               <dbl>        <dbl>
## 1 TRUE             6300 Switzerland            NA        0.556
## 2 TRUE             6300 Switzerland            NA        0.556
## 3 TRUE             6300 Switzerland            NA        0.556
```

Exercises `mutate`

- What's the percentage of people that died out of those that got injured?

Exercises `mutate`

- What's the percentage of people that died out of those that got injured?

- Out of those that died, how many got injured before?

Exercises `mutate`

▶ What's the percentage of people that died out of those that got injured?

▶ Out of those that died, how many got injured before?

▶ Create a new column that is the concatenation of the member id and it's citizenship.

summarise

summarise

It's just like mutate, but the output has to be of only one row, and it eliminates (in the output) everything else that is not mentioned*:

summarise

It's just like mutate, but the output has to be of only one row, and it eliminates (in the output) everything else that is not mentioned*:

```
members1905 %>%
  summarise(anyone_died = any(died),
          max_height = max(highpoint_metres,na.rm = TRUE),
          last_citizenship = last(citizenship),
          youngest = min(age),
          percent_dead = mean(died))
```

```
## # A tibble: 1 x 5
##   anyone_died max_height last_citizenship youngest percent_dead
##   <lgl>           <dbl> <chr>               <dbl>        <dbl>
## 1 TRUE             6300 Switzerland            NA        0.556
```

We are only getting one row, and losing everything else. Using, filter, mutate and
summarise, indicate how far the average of the highpoint metres for the women from
France differs from the average from all the table (use na.rm=TRUE).

group_by

`group_by`

Figure 5: Another type of group

group_by

```
members %>%
  group_by(died) %>%
  summarise(Percentage_injured = mean(injured))
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 2 x 2
##   died  Percentage_injured
##   <lgl>              <dbl>
## 1 FALSE            0.0227
## 2 TRUE             0.00271
```

What do you think group_by does?

group_by

On it's own, `group_by` doesn't do much, it really shines when we combine it with the other 5 verbs that we have studied.

group_by

On it's own, group_by doesn't do much, it really shines when we combine it with the other 5 verbs that we have studied.

All it does is group by the variables we tell it to, and the following modifications that happen on the data.frame are done on each of the groups we defined, as if each group was a dataframe.
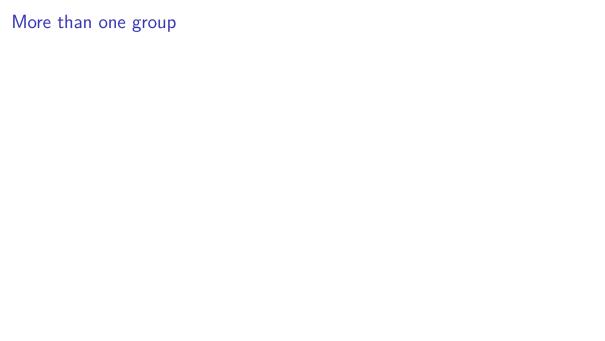
group_by + mutate

When we combine it with mutate, we get a new observation, that is only different for each of the grouped variables.

```
members_n %>% select(-(4:14)) %>%
  group_by(season) %>%
  mutate(Obs_per_season = n())
```

```
## # A tibble: 1,106 x 4
## # Groups:   season [4]
##    peak_name    year season Obs_per_season
##    <chr>       <dbl> <chr>           <int>
##  1 Ama Dablam   1979 Autumn            493
##  2 Ama Dablam   1983 Autumn            493
##  3 Ama Dablam   1983 Autumn            493
##  4 Ama Dablam   1985 Spring            555
##  5 Ama Dablam   1988 Spring            555
##  6 Ama Dablam   1992 Spring            555
```

group_by + mutate

Repeat what we did above, but grouping by year, instead of season.

More than one group

More than one group

If you want to group with more than one variable, you can simply add the column in the argument:

```
members %>%
  group_by(year,season) %>%
  summarise(Average_height = mean(highpoint_metres,
                                  na.rm = TRUE))
```

```
## `summarise()` regrouping output by 'year' (override with `.groups` argu

## # A tibble: 239 x 3
## # Groups:   year [92]
##     year season Average_height
##    <dbl> <chr>          <dbl>
## 1  1905 Summer          6300
## 2  1907 Autumn          7270
## 3  1909 Autumn          6965
## 4  1910 Spring          6965
```
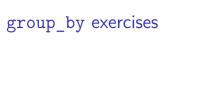
group_by exercises

group_by exercises

- What year had the most dead people? (Hint: arrange)

group_by exercises

▶ What year had the most dead people? (Hint: arrange)

▶ Using only those with a role of climbers, compute the mean and standard deviation of the yearly number of climbers for each season. (Hint: use season as the first grouping variable)

group_by exercises

▶ What year had the most dead people? (Hint: arrange)

▶ Using only those with a role of climbers, compute the mean and standard deviation of the yearly number of climbers for each season. (Hint: use season as the first grouping variable)

▶

group_by + summarize

Answer to the second question above:

```r
members %>%
  filter(expedition_role =='Climber') %>%
  group_by(season,year) %>%
  summarize(Total = n()) %>%
  summarise(Mean = mean(Total),
            Std = sd(Total))

## `summarise()` regrouping output by 'season' (override with `.groups` arg

## `summarise()` ungrouping output (override with `.groups` argument)

## # A tibble: 5 x 3
##   season  Mean   Std
##   <chr>   <dbl> <dbl>
## 1 Autumn  321.   270.
## 2 Spring  256.   240.
```

group_by + filter

If you want to eliminated the groups that don't have enough, or that have too many observations, you can do it by combining filter and group_by directly:

```
(surviving_members <- members %>%
  group_by(expedition_id) %>%
  filter(n() > 10))
```

```
## # A tibble: 38,762 x 21
## # Groups:   expedition_id [2,264]
##    expedition_id member_id peak_id peak_name  year season sex     age
##    <chr>         <chr>     <chr>   <chr>     <dbl> <chr>  <chr> <dbl>
##  1 AMAD79101     AMAD7910~ AMAD    Ama Dabl~  1979 Spring M        35
##  2 AMAD79101     AMAD7910~ AMAD    Ama Dabl~  1979 Spring M        37
##  3 AMAD79101     AMAD7910~ AMAD    Ama Dabl~  1979 Spring M        23
##  4 AMAD79101     AMAD7910~ AMAD    Ama Dabl~  1979 Spring M        44
##  5 AMAD79101     AMAD7910~ AMAD    Ama Dabl~  1979 Spring M        25
##  6 AMAD79101     AMAD7910~ AMAD    Ama Dabl~  1979 Spring M        28
```

`rowwise`

This function, as it name tells us, is like doing a `group_by`, but operates on each row. This one is particular useful when you are creating your own functions and they have *weird* interactions with vectors, like using `sum,` `mean,` and such... But we would probably get different results on each row.

Thank you :)

Any questions?

Further references

- Check the help page and the vignettes of dplyr! (type ?dplyr, or: `vignette('dplyr')` on the console)
- R for Data Science, by Hadley Wickham
- Advanced R, by Hadley Wickham
- The R Inferno, by Patrick Burns